

Cayman Roden

QA Automation Engineer - 3,500+ Tests Across Production Repos - pytest - Playwright - Pact - Locust - CI/CD - LLM Evaluation

310-982-0492 | caymanroden@gmail.com | [LinkedIn](#) | [GitHub](#) | [Portfolio](#) | Remote, US-based

See it running: [DocExtract](#) | [EnterpriseHub](#) | [PyPI: mcp-server-toolkit](#)

Open source: [LiteLLM PR #24551](#) (47K+ stars, open / under review) | [FastAPI PR #15217](#) (98K+ stars, open / under review)

Professional Summary

QA Automation Engineer with 3,500+ automated tests shipped across production repositories (docextract 1,280 at 81% coverage, mcp-server-toolkit 600 at ~83%, jorge-real-estate-bots 1,700+), all developed TDD-first. Designed test automation frameworks using pytest, Playwright, Pact, Locust, and GitHub Actions. Built LLM evaluation pipelines with golden and adversarial test sets, rubric-based judge scoring, LLM-as-judge CI gates, faithfulness checks, and an offline replay job that gates merges deterministically at zero API cost. Full-stack test ownership from framework selection through coverage strategy, production validation, and handoff.

Technical Skills

Test Frameworks: pytest | pytest-asyncio | Playwright | Selenium | Vitest | unittest | Postman | Pact | Locust

Test Types: Unit | Integration | E2E | API Contract | Consumer-Driven Contract | Load | Regression | Adversarial | LLM Evaluation | RAG Evaluation | Faithfulness | Accessibility (WCAG)

CI/CD and Quality: GitHub Actions | Jenkins | Docker | ARQ | Render | Vercel | Cloud Run | AWS | bandit | pip-audit | mypy strict | ruff | pre-commit | Coverage Enforcement | OpenTelemetry | Prometheus | Grafana

Languages and Systems: Python | TypeScript | JavaScript | SQL | FastAPI | PostgreSQL | Redis | REST API | WebSockets | Async | GoHighLevel Webhooks | Claude API | OpenAI | LiteLLM | RAG | Embeddings | Vector Databases | pgvector | LangGraph | MCP | Tool Calling | Golden Sets | Regression Suites

Experience

QA Automation Engineer | Acuity Real Estate (Contract) | Remote | January 2026 - March 2026

Sole QA owner for a production AI SMS bot platform serving a real estate client. Established testing standards for a greenfield codebase from day one.

- Set test strategy before production code merged: framework selection (pytest-asyncio), coverage thresholds, CI gate configuration, and branch protection rules.
- Validated client-facing behavior across the full conversation and CRM flow: 500+ real leads processed with no routing errors, unhandled exceptions, or client-visible failures across the 3-month deployment.
- Delivered 1,700+ tests at handoff with a formal 11-section test-strategy document, coverage report, runbook, and risk-based prioritization for new conversation flows.
- Zero production regressions and zero reported incidents across the deployment lifecycle.

QA Automation Engineer | Contract • Remote | June 2022 - Present

- Shipped 3,500+ automated tests across production repos using TDD-first methodology: failing test, minimal implementation, refactor, verify.
- Built CI/CD pipelines with GitHub Actions including SAST scanning, dependency CVE auditing, type checking, ruff linting, and coverage enforcement gates.
- Designed LLM evaluation infrastructure: golden test sets, adversarial fixtures, rubric-based judge scoring, faithfulness checks, RAG evaluation, and an offline CI replay job that runs the eval gate at zero API cost.
- Published [mcp-server-toolkit](#) to PyPI with a custom MCPTestClient enabling unit testing of MCP servers without live API keys; 600 tests at ~83% coverage.
- Authored [LiteLLM PR #24551](#): typed exception mapping for Router fallback so distinct error types surface; tests for 4 error cases.

Test Automation Portfolio

[EnterpriseHub](#) - Multi-Agent Orchestration Platform

pytest | pytest-asyncio | mypy strict | bandit | pip-audit | GitHub Actions

- Designed test suite covering the 22-agent mesh, inter-agent handoff verification, per-agent model-routing validation, and autonomous task-delegation correctness.
- Tested L1 memory / L2 Redis / L3 PostgreSQL cache behavior: hit/miss/eviction edge cases, TTL behavior, and eviction-policy correctness.
- Security testing: Ed25519 webhook signature verification, parameterized SQL injection resistance, rate limiting enforcement, and OWASP control validation.

[DocExtract](#) - LLM Evaluation Suite

pytest | RAG evaluation | LLM-as-judge | Brier score calibration

- Designed an eval corpus with golden and adversarial cases; headline metric is **95.5% accepted F1 on a 28-case golden baseline** (replayed deterministically in CI at zero API cost), backed by a full 72-case corpus (51 golden + 21 adversarial).
- Built an evaluation pipeline covering context recall, faithfulness, answer relevancy, evidence extraction, and CI regression gates.
- Tracked confidence calibration with Brier score checks across fixture types.

[Jorge Real Estate AI](#) - Production SMS Bot Suite

pytest | pytest-asyncio | Pact | Locust | GitHub Actions

- Designed a multi-layer test strategy for a greenfield codebase: async event-loop correctness, parameterized fixtures, all routing branches across 3 bots, and bilingual EN/ES conversation paths.
- Tested dual-scheme webhook signature verification (HMAC + RSA), replay prevention, malformed payload rejection, rate limiting, atomic session locks, and state recovery.
- Load tested against production: p50 130ms, p95 250ms, p99 400ms at 5.54 req/s with zero 5xx errors.

Education and Continuous Learning

11 professional certificates, ~750 hours (Google, Microsoft, IBM, Duke, Vanderbilt, Linux Foundation, DeepLearning.AI), including Duke LLMOps, Vanderbilt Prompt Engineering, IBM Generative AI Fundamentals, Google Cloud Generative AI Leader, and Linux Foundation OSS / Linux / Git.