

Cayman Roden

AI Solutions Engineer - Production LLM Systems - API Integrations - CRM Automation

310-982-0492 | caymanroden@gmail.com | [LinkedIn](#) | [GitHub](#) | [Portfolio](#) | Remote, US-based

Proof links: [EnterpriseHub demo](#) | [DocExtract demo](#) | [PyPI: mcp-server-toolkit](#)

Open source: [LiteLLM PR #24551](#) (47K+ stars, open / under review) | [FastAPI PR #15217](#) (98K+ stars, open / under review)

Professional Summary

AI Solutions Engineer: 4 years shipping production AI systems on top of 10 years of client-facing delivery. Scoped, architected, and shipped Acuity Real Estate's bilingual lead-qualification platform end to end: 500+ leads, zero reported incidents over a 3-month run, 1,700+ tests at handoff. Turns ambiguous client workflows into deployed AI systems with Claude API integration, GoHighLevel CRM automation, FastAPI services, and proof-of-concept-to-production discipline. Comfortable owning the customer relationship from discovery through technical demo, integration, and handoff.

Technical Skills

Python | FastAPI | REST APIs | OpenAPI | PostgreSQL | Redis | Docker | ARQ | Render | Vercel | Cloud Run | AWS | GitHub Actions CI/CD | API Integration | Solution Architecture | SaaS Architecture | Claude API | OpenAI | RAG | Embeddings | Vector Databases | pgvector | Agentic AI | MCP | LangGraph | Tool Calling | LiteLLM | Prompt Engineering | LLM Evals | Golden Sets | Regression Suites | OpenTelemetry | Cost Attribution | GoHighLevel | Webhooks | Proof of Concept | Stakeholder Management | Customer Discovery | Technical Demos | Customer Success | Async | WebSockets

Recent Production Engagement

Solutions Engineer and AI Architect | Acuity Real Estate (Contract) | Remote | Jan 2026 - Mar 2026

- Ran technical discovery with the client, designed the GoHighLevel CRM webhook architecture for bilingual lead qualification, and selected the Claude API on NLP quality and cost.
- Built 3 specialized SMS bots (Lead Intake, Buyer Qualification, Seller Qualification) with cross-bot routing, bilingual EN/ES support, and an admin handoff protocol.
- Production results: 500+ leads at sub-500ms p95, zero reported incidents over the 3-month run, 226 GHM workflows audited, 12 bugs diagnosed in one alignment session. Delivered 1,700+ automated tests at handoff (conversation flows, webhook contracts, rate limiting, recovery paths) with zero production regressions.
- Delivered [DocExtract](#) as an API-first document-extraction service: 95.5% accepted F1 on a 28-case golden baseline (CI-replayed at zero API cost), a 72-case corpus (51 golden + 21 adversarial), and an LLM-as-judge CI gate; ~\$0.03/doc and ~4.1s p95 are modeled from an in-repo pricing table.

Experience

Solutions Engineer | Contract · Remote | June 2022 - Present

- Built [EnterpriseHub](#): an AI operations platform with 7-agent orchestration (~10 with auto-discovery), 4-dimension routing (success/load/cost/latency), prompt versioning, cost controls, eval gates, and a live Streamlit demo for stakeholder review.
- Published [mcp-server-toolkit](#) to PyPI: production MCP servers with JWT/OAuth 2.1 auth, OpenTelemetry tracing, per-span cost attribution, a cache layer with a measured 2.9x speedup, and 600 tests at ~83% coverage.
- Built a multi-tenant embeddable chat widget (FastAPI, pgvector, Claude API, Next.js, TypeScript, Stripe billing, per-tenant isolation, usage analytics, audit logging) for SaaS-style client pilots.
- Delivered technical documentation, BRDs, runbooks, deployment guides, and handoff materials across AI, API, analytics, and web engagements.

Education and Continuous Learning

11 professional certificates, ~750 hours (Google, Microsoft, IBM, Duke, Vanderbilt, Linux Foundation, DeepLearning.AI), including Duke LLM Ops, Vanderbilt Prompt Engineering, IBM Generative AI Fundamentals, Google Cloud Generative AI Leader, and Linux Foundation OSS / Linux / Git.